SEEING FLOW

Video Cloud Solution

Jun 2026 Edition

SEEING FLOW

Technical White Paper : Video Cloud Solution	
Version No.	Ver. 1.2
Date Prepared.	16 Jun. 2026
Prepared By	Seeingflow
Reviewed By	Seeingflow
63 Pages Total	

Disclaimer: All information presented in this white paper is considered the intellectual property and proprietary asset of Seeingflow. The content within is provided solely for informational purposes, and while every effort has been made to ensure its accuracy, Seeingflow makes no representations or warranties of any kind, express or implied, regarding the completeness, accuracy, reliability, suitability, or availability of the information, products, or services discussed herein. Any reliance you place on such information is therefore strictly at your own risk.

Video Cloud Solution
Table of Contents

Video Cloud Solution

Executive Summary

The Video Cloud is a reusable cloud media platform for ingesting video, storing recordings, processing media, distributing playback, exposing video services, and managing operations. It is the technical foundation on which camera services, enterprise video applications, advanced analytics, and future commercial products can be built. The platform is designed for private-cloud deployment, RTSP/ONVIF camera integration, H.264/H.265 video handling, mobile-first playback, web-based administration, and scalable operation toward millions of connected cameras.

The recommended solution emphasizes enterprise-grade security, privacy protection, service observability, and user experience. Video content is inherently sensitive, so the platform must treat access control, encryption, retention enforcement, secure deletion, audit trails, and compliance evidence as core capabilities. At the same time, users should experience fast live view, reliable recording search, clear service states, and consistent playback across mobile and web terminals.

SEEING FLOW

Seeingflow Technology Pte. Ltd.
100 Lorong 23 Geylang
#02-02, D'Centennial Building,
Singapore 388398

Video Cloud Solution

1. Introduction

A Video Cloud solution is the core platform layer for receiving video streams and uploaded media, storing video records, processing media assets, distributing playback to user terminals, and managing users, permissions, usage records, monitoring, and operations. It is designed as a reusable foundation that can support multiple business scenarios without being tied to a specific network provider, device brand, analytics product, or commercial package.

The solution supports IP camera recording, live streaming, playback, video-on-demand, snapshot capture, watermarking, transcoding, and secure content delivery. It can serve as the foundation for surveillance services, consumer camera services, smart-space platforms, enterprise video management, live event recording, and media applications. The target architecture should be capable of growing toward millions of connected cameras while maintaining consistent service quality.

This paper focuses on the Video Cloud platform itself. Hardware procurement, network-service bundles, downstream analytics modules, and user-facing commercial packaging may integrate with the platform, but they are not part of the core scope. The objective is to define a professional, scalable, secure, and operationally mature Video Cloud foundation.

2. Design Assumptions

2.1 Platform Scale & Architecture

The Video Cloud should be planned as a large-scale platform rather than a project-specific media server. The long-term design target is millions of connected cameras, which means the system must avoid single points of capacity in the access gateway, metadata database, storage namespace, stream-processing queue, and monitoring pipeline. Horizontal scaling, partitioned data models, distributed object storage, asynchronous processing, and device/channel-level observability should be treated as baseline architectural principles.

Video Cloud Solution

2.2 Recording Retention

The Video Cloud should be planned as a large-scale platform rather than a project-specific media server. The long-term design target is millions of connected cameras, which means the system must avoid single points of capacity in the access gateway, metadata database, storage namespace, stream-processing queue, and monitoring pipeline. Horizontal scaling, partitioned data models, distributed object storage, asynchronous processing, and device/channel-level observability should be treated as baseline architectural principles.

2.3 Client Terminals & Playback Experience

Most playback is expected to occur through a mobile application, as users commonly access video while away from the monitored location. A web portal should additionally be provided for browser-based playback, administration, customer support, and operational supervision.
This dual-terminal model has direct implications for authentication, playback authorization, session management, UI design, and testing strategy. The mobile experience must be optimized for speed and simplicity; the web portal must support more structured management workflows.

2.4 Device Integration & Camera Compatibility

The primary integration model is RTSP and ONVIF, with H.264 and H.265 as the supported video codecs. While these standards are widely adopted in IP camera deployments, real-world interoperability still requires systematic compatibility testing across camera vendors, firmware versions, stream profiles, bitrate settings, keyframe intervals, and network conditions.

The access layer should normalize device behavior as much as possible and expose stable internal stream states to the rest of the platform.

Video Cloud Solution

2.5 Private Cloud Deployment

Deployment is expected on a private cloud infrastructure. This places stronger emphasis on private network segmentation, controlled ingress and egress, internal service discovery, private container or VM clusters, dedicated storage pools, private monitoring systems, and governed access from external playback clients. The solution should adopt market-standard cloud-native patterns while operating independently of public-cloud managed services.

2.6 Peak Load & Capacity Planning

Playback demand is not uniform across the day. Peak usage is expected before 9:00 AM on workdays, during lunch hours, after 6:00 PM, and throughout weekends. Capacity planning must therefore model real human behavior rather than average daily traffic.

Access gateways, authorization services, CDN and origin capacity, metadata query paths, and mobile playback APIs should all be stress-tested against these concentrated usage windows.

2.7 Storage & CDN Architecture

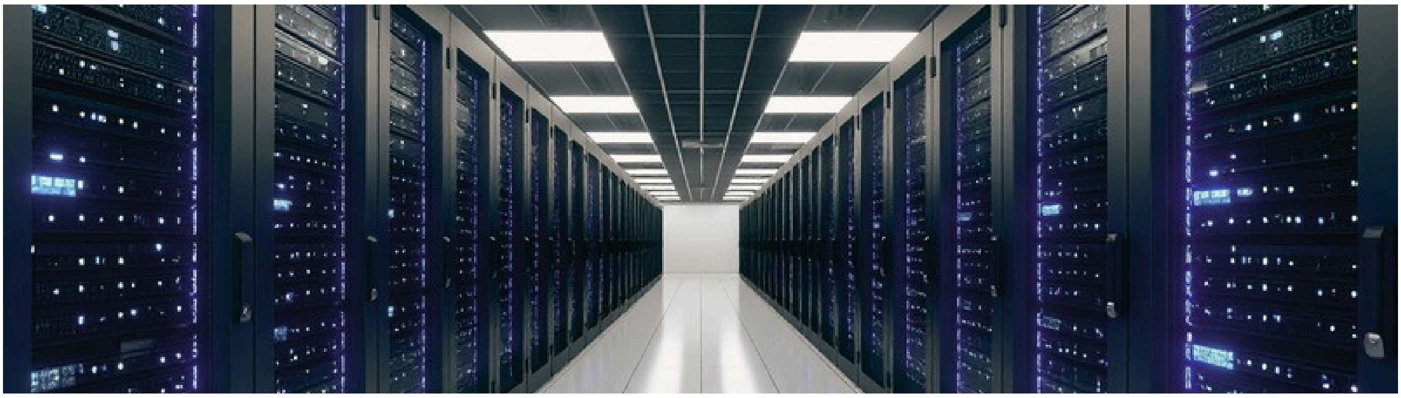

Storage and CDN design should follow market-standard patterns and mature engineering practices. This includes distributed object storage with replication or erasure coding, lifecycle management, media segment organization, HTTP-based delivery, edge caching, origin protection, TLS, signed playback access, and cache-control policies.

Where required, regional delivery nodes should be provisioned. Clear separation must be maintained between hot playback data, retained recordings, archive material, logs, and operational metadata.

3. Scalability and Private Cloud Strategy

At million-camera scale, the Video Cloud must be engineered as a distributed service fabric. Cameras should not connect to a single central server group. Instead, access gateways should be deployed in scalable pools, and device connections should be routed through load balancing, regional ingress, or edge access domains. Each gateway should be stateless or minimally stateful where possible, so capacity can be added without complicated migration.

The metadata architecture is as important as the media architecture. Every recording, camera, stream, retention policy, user permission, playback session, and operational event creates metadata. If metadata is not partitioned carefully, playback and recording queries can become the limiting factor even when storage capacity is sufficient. The platform should therefore define sharding or partitioning strategies early, including partition keys such as tenant, region, device/channel, time window, and recording type.
The private-cloud environment should be designed with clear separation of duties. Public-facing gateways, internal media services, storage systems, databases, monitoring tools, administrative systems, and security services should run in segmented network zones. Administrative access should be restricted through controlled bastion or privileged-access workflows, and production data paths should be separated from development, testing, and analytics environments.

Capacity management should be driven by service indicators rather than infrastructure size alone. Important indicators include active camera count, online camera ratio, live stream ingest rate, recording write throughput, playback session concurrency, authorization request rate, storage object count, CDN cache hit ratio, metadata query latency, and error rate by camera model or firmware version. These measurements help the platform team expand capacity before user experience is affected.

Video Cloud Solution

4. Value and Benefits

Centralized Video Service Foundation

The Video Cloud provides one common backend for live video ingest, recording, storage, playback, processing, and management, reducing duplicated platform development across projects.

Scalable Storage and Playback

Distributed storage and CDN delivery allow the system to support growing camera counts, larger video libraries, and distributed viewers without redesigning the platform.

Operational Efficiency

Centralized monitoring, permission management, logs, lifecycle policies, and automated workflows reduce manual maintenance and make incidents easier to diagnose.

Better User Experience

Cloud playback, adaptive streaming, CDN acceleration, and terminal-friendly protocols improve startup speed, stability, and playback quality across mobile, web, and desktop terminals.

Security and Compliance Readiness

Access control, private storage, encrypted transmission, audit logs, watermarking, retention policy, and tenant isolation help the platform serve enterprise and regulated environments.

Future Extensibility

Once the Video Cloud is stable, other services such as advanced analytics, smart alerts, service packages, client applications, or partner ecosystems can be added through APIs without rebuilding the media foundation.

Enterprise-Grade Trust

A high-end Video Cloud should not only process video; it should provide assurance that recordings are protected, access is controlled, operations are auditable, and service behavior remains predictable even as the platform grows. This trust layer is one of the most important differences between a simple video server and a serious cloud video platform.

5. User Experience Principles

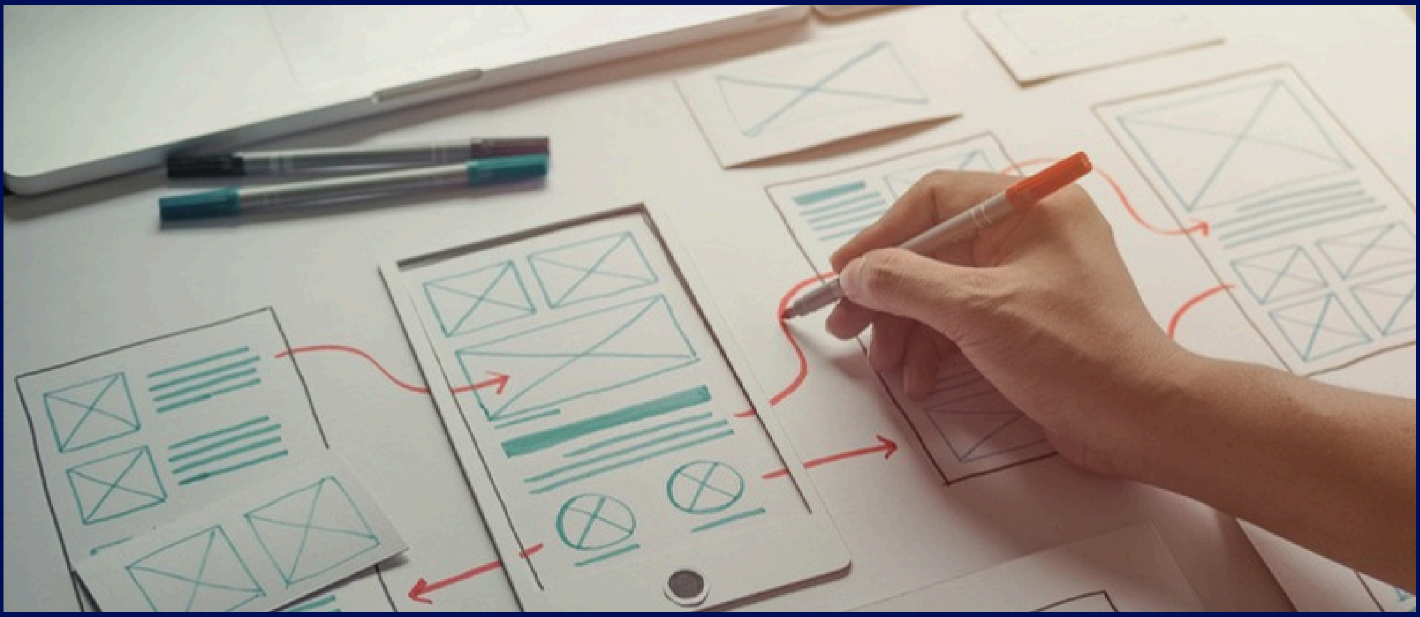

User experience is a core technical requirement, not a cosmetic layer. Users will judge the platform by how quickly they can open the mobile app, confirm a camera is online, start live playback, locate a recent recording, and understand the root cause of any playback issue — whether it originates from the camera, network, account permissions, or the platform itself. The platform must therefore expose clear, actionable service states rather than vague error messages.

5.1 Mobile Experience

Mobile playback should be optimized for short, frequent sessions. Users typically access the app before work, during lunch, after work, or on weekends to quickly check a location.

The experience must prioritize:

- **Fast Startup and Stable Reconnect** — minimize time from app launch to live view
- **Adaptive Playback** — adjust stream quality gracefully under weak or variable network conditions
- **Low Battery Impact** — efficient background behavior and stream handling
- **Timeline Preview** — quick access to recent recordings without deep navigation
- **Clear Loading and Error States** — users should always know what is happening and why

A technically powerful platform will still feel poor if playback startup is slow or error handling is confusing.

5.2 Web Portal Experience

The web portal should support more deliberate, structured workflows. Administrators and support teams require device search, account lookup, timeline review, retention status, event history, audit records, role management, and system-health visibility. The portal should be designed for operational efficiency and control, while the mobile app is designed for speed and confidence.

5. User Experience Principles

5.3 Experience Metrics

The platform should measure user experience continuously. Recommended metrics include:

- **App startup time and live-view startup time**
- **First-frame time and playback failure rate**
- **Rebuffering ratio and timeline query latency**
- **Recording availability and camera offline duration**
- **Support-ticket correlation against platform and stream events**

These metrics make experience quality visible across engineering, operations, and management — enabling data-driven decisions rather than reactive fixes.

5.4 Complexity Abstraction

A premium Video Cloud protects users from underlying complexity. Protocol translation, stream normalization, retention policy enforcement, CDN and origin selection, and security token exchange should all occur behind the scenes. Users should experience a simple, reliable service — while the platform silently absorbs the complexity required to deliver it.

6. Architecture

The recommended Video Cloud architecture is organised across six independently scalable and observable layers: access, storage, processing, distribution, application, and management.

Figure 1. Six-Layer Architecture of Video Cloud

6. Architecture

6.1 Access Layer

Terminal devices, cameras, upload clients, and playback clients connect through access servers and gateway services. This layer is responsible for stream ingest, upload and playback requests, authentication handoff, rate limiting, connection management, and protocol adaptation.

For RTSP and ONVIF camera environments, the access layer additionally manages camera registration, stream status monitoring, protocol compatibility, reconnect behavior, and normalization of differing vendor implementations.

6.2 Storage Layer

Video files, segments, snapshots, thumbnails, manifests, and metadata are persisted across distributed object and file storage systems with supporting databases. The storage layer must support:

- **Retention tiers** — three-day baseline with optional extension up to 30 days
- **Lifecycle policies** — automated hot, warm, and cold tier transitions
- **Redundancy and encryption** — at rest and in transit
- **Secure deletion** — policy-enforced, audit-traceable removal

At large scale, storage design must account for object count, write amplification, timeline lookup speed, metadata consistency, and secure deletion behavior — not capacity alone.

6.3 Processing Layer

Stored and live media can be transcoded, packaged, edited, watermarked, clipped, indexed, and converted into formats suitable for different devices and network conditions. Screenshot capture and preview generation also reside in this layer.

Given H.264 and H.265 as the primary codecs, the processing layer must support codec-aware handling, profile validation, adaptive output policy, and efficient processing queues that do not block live ingest or playback paths.

6.3 Processing Layer

Stored and live media can be transcoded, packaged, edited, watermarked, clipped, indexed, and converted into formats suitable for different devices and network conditions. Screenshot capture and preview generation also reside in this layer.

Given H.264 and H.265 as the primary codecs, the processing layer must support codec-aware handling, profile validation, adaptive output policy, and efficient processing queues that do not block live ingest or playback paths.

6. Architecture

6.4 Distribution Layer

Processed video is delivered through CDN nodes or edge caches to improve playback latency, reduce origin load, and support geographically distributed users. In a private-cloud deployment, this layer combines private origin services, edge caching, controlled internet gateways, and market-standard HTTP media delivery patterns to sustain stable playback during peak viewing windows.

6.5 Application Layer

The application layer exposes live streaming, video-on-demand, playback session control, recording query, video analytics hooks, API services, and integration points for mobile apps, web portals, enterprise systems, and partner applications.

This layer abstracts media complexity from client applications and provides consistent APIs for live view, playback authorization, recording timeline search, device and channel management, and usage records.

6.6 Management Layer

The management layer governs user accounts, tenant isolation, permissions, usage records, system configuration, device and channel registration, monitoring, alerting, logs, audit trails, service health, and operations dashboards.

In an enterprise platform, management functions are not secondary tooling — they are the control plane for operating millions of devices, enforcing retention policy, demonstrating compliance, and maintaining service quality at scale.

6.7 Architectural Summary

Layer	Main Function	Typical Components
Access	Receive upload/playback requests and live streams	Access gateway, stream ingest, API gateway, auth handoff, rate limiting
Storage	Store video files, segments, snapshots, and metadata	Object storage, distributed file system, metadata database, lifecycle policy
Processing	Transcode, package, screenshot, clip, watermark, and index video	Transcoder, job queue, media workers, workflow state
Distribution	Deliver video to terminals efficiently	CDN, edge cache, origin protection, HTTPS delivery
Application	Expose video services to user-facing systems	Live streaming API, VOD API, recording query, playback authorization
Management	Operate and govern the platform	User management, permission management, usage records, monitoring, audit logs

The six layers operate as a coordinated system rather than isolated components. Access receives and normalises traffic; storage preserves media and metadata; processing prepares usable outputs; distribution accelerates delivery; the application layer exposes video functions to clients; and management governs the entire system. This logical separation ensures each layer can scale, fail, and recover independently without cascading impact across the platform.

7. Reference Workflows

The workflow diagrams below translate the six-layer architecture into the main operating scenarios of the platform. They are intended to show how responsibilities move across the system, not to prescribe a single vendor implementation.

7.1 Live Stream Workflow

This diagram explains the low-latency path for live viewing. It emphasizes that ingest stability, session authorization, delivery performance, and monitoring must work together in real time.

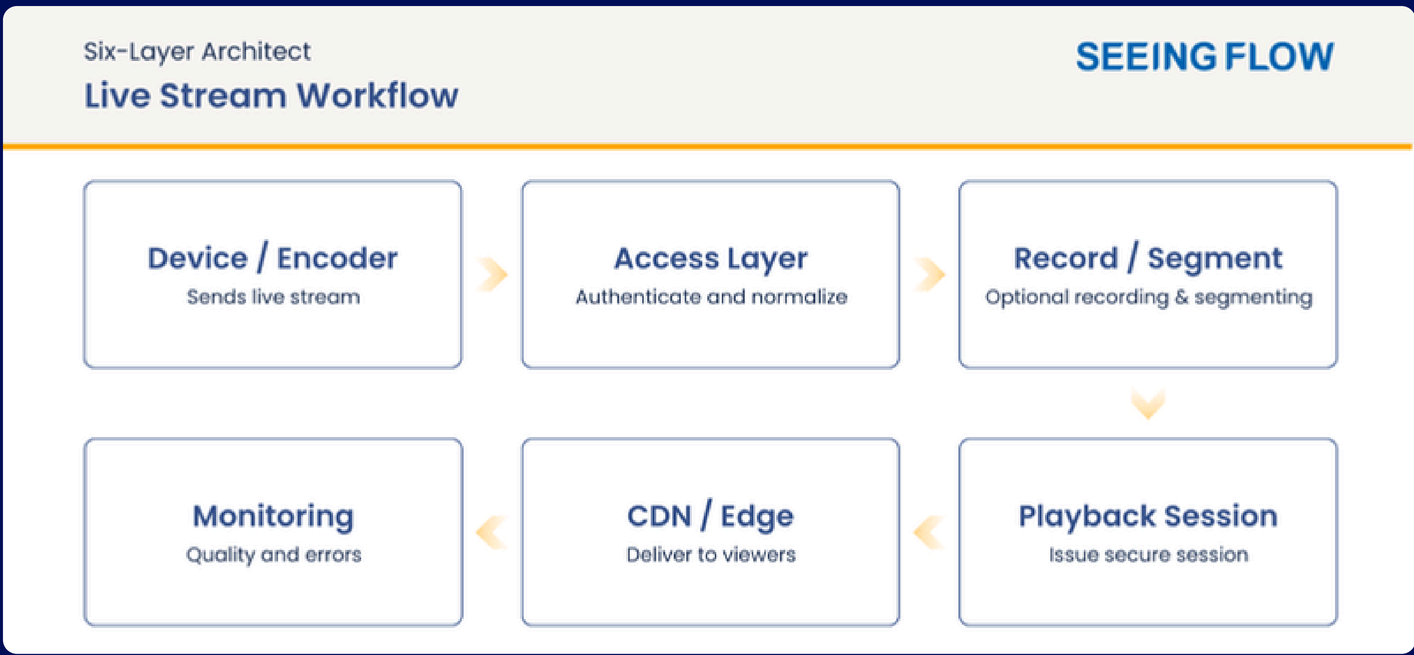

Figure 2. Live stream workflow

7.2 Recording and Playback Workflow

This diagram explains how incoming streams become searchable and playable recordings. The core design concern is timeline indexing, retention enforcement, and secure retrieval.

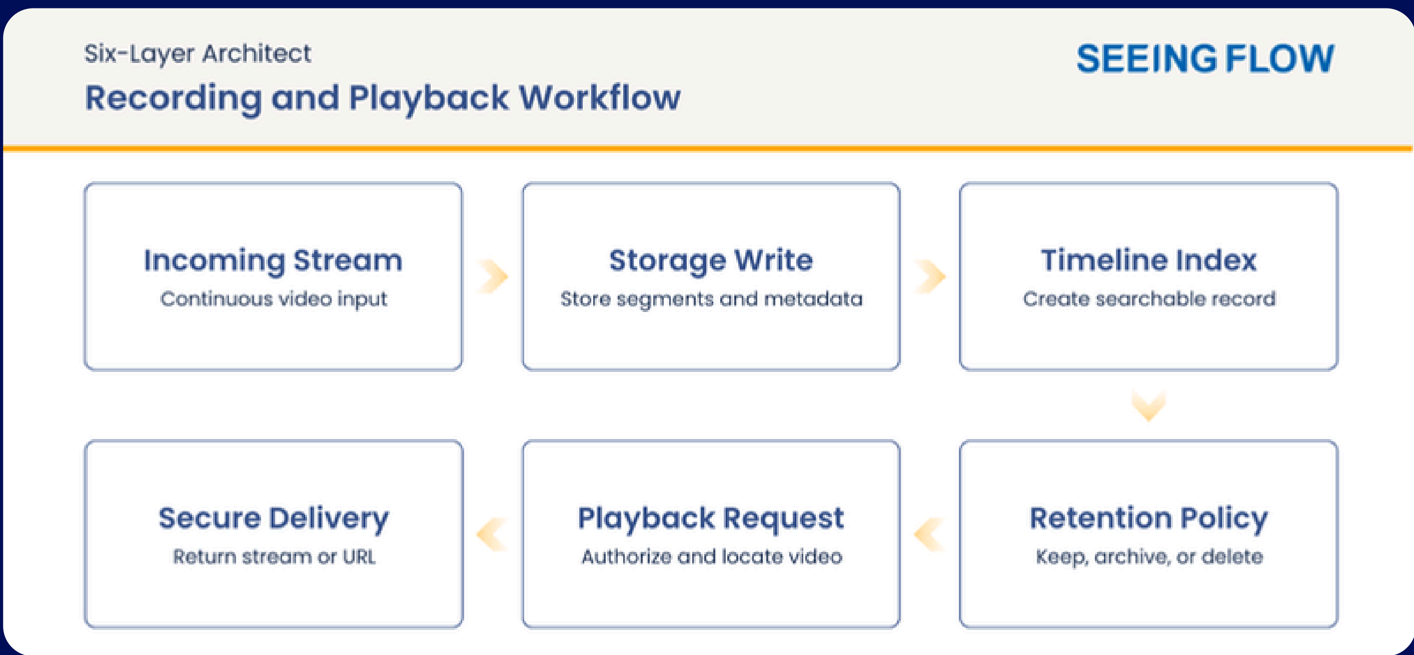

Figure 3. Recording and playback workflow

7. Reference Workflows

7.3 Upload / VOD Processing Workflow

This diagram explains batch processing for uploaded media. The workflow is less latency-sensitive than live streaming, but it requires strong validation, repeatable processing, and consistent output packaging.

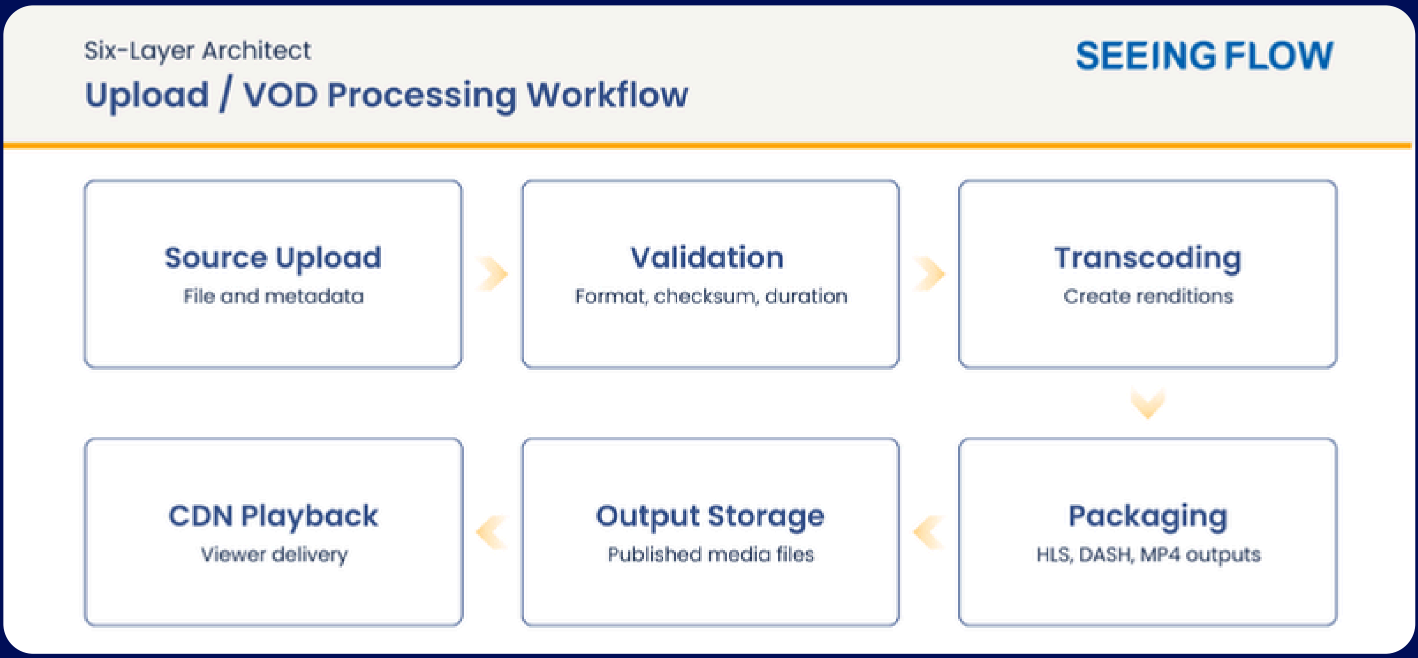

Figure 4. Upload and VOD processing workflow

7.3 Management and Operations Workflow

This diagram explains how the platform is operated after launch. It connects telemetry, dashboards, alerts, reporting, and optimization so administrators can manage quality, capacity, and reliability.

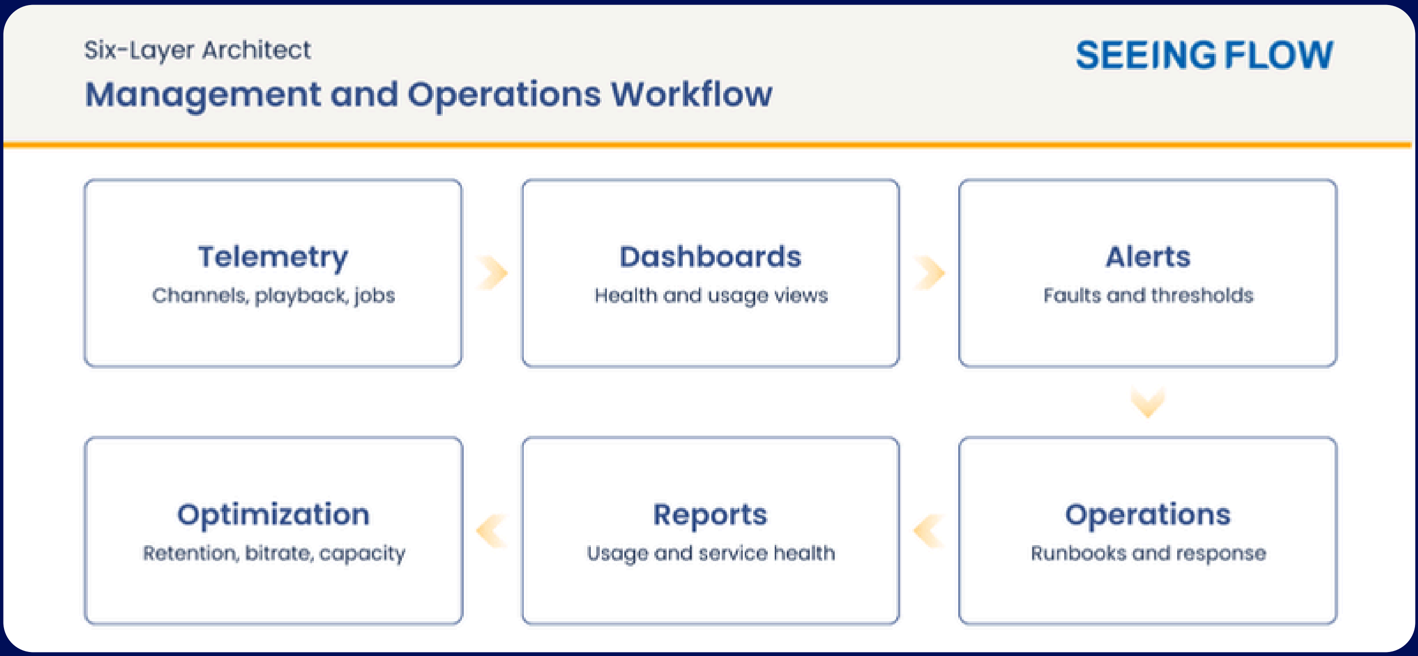

Figure 5. Management and operations workflow

8. Implementation Approach

Build the Video Cloud as a modular platform

Keep access, storage, processing, distribution, application APIs, and management services separated so each can scale and evolve independently. This modularity is important for commercial flexibility as well as engineering quality: different customers may have different scale, retention, security, and integration requirements, but the platform foundation should remain consistent.

Start with core video functions

The first release should include stream ingest, recording, playback, storage retention, user/permission management, monitoring, and basic processing such as transcoding, snapshots, and watermarking. These capabilities define whether the platform is dependable. Advanced features should not distract from the reliability of the live-view, recording, and playback foundation.

Use standard media protocols and formats

Support RTSP and ONVIF as primary camera integration protocols, H.264 and H.265 as primary codecs, HLS or DASH for scalable playback, MP4 for file output where required, and HTTPS delivery for terminal playback.

Design for multi-tenant operation

Even if the first deployment has one business customer, the platform should support tenant IDs, separated storage namespaces, role-based access, audit logs, and per-tenant usage accounting.

Optimize storage policy early

Video storage is a major operating factor. Define recording retention, hot/warm/cold tiers, automatic deletion, archive policy, and separate source/output/log storage from the beginning. The three-day baseline and optional 30-day retention model should be encoded as policy, so different service packages can be activated without manual storage intervention.

Instrument service quality

Track stream ingest success, camera/channel online rate, playback startup time, buffering, CDN cache hit ratio, processing duration, job failure rate, storage growth, and platform utilization.

Expose clean APIs

Client applications, analytics services, service platforms, and partner portals should integrate through documented APIs for device/channel registration, playback authorization, recording query, file upload, processing jobs, user permissions, and usage records.

Plan acceptance testing around realistic user behavior

Test cases should include morning, lunch, evening, and weekend playback peaks; unstable camera networks; mobile network switching; long timeline queries; mass camera reconnect; retention expiration; and administrator audit review. This creates a more credible launch process than simple functional testing.

Video Cloud Solution

9. Security, Privacy, and Compliance Framework

Security and privacy are defining characteristics of the Video Cloud, not afterthoughts. The platform handles sensitive video content, account data, device metadata, location records, playback history, administrative activity, and operational logs. Regardless of how the platform is marketed, the nature of video data makes confidentiality, integrity, availability, privacy, and auditability central to the product promise.

ISO/IEC 27001:2022

Client applications, analytics services, service platforms, and partner portals should integrate through documented APIs for device/channel registration, playback authorization, recording query, file upload, processing jobs, user permissions, and usage records.

ISO/IEC 27017:2015

ISO/IEC 27017:2015 provides cloud-specific control guidance applicable even in private-cloud environments, covering virtual environment separation, shared responsibility boundaries, operational procedures, monitoring, and network alignment. The platform should formally document which controls belong to the platform operator, infrastructure teams, and integrated application teams respectively.

ISO/IEC 27701

Where personal data is involved, ISO/IEC 27701 governs privacy management. Video recordings, camera metadata, account information, device location, access logs, and playback history are all potentially privacy-sensitive. Privacy roles, data classification, consent records, retention rules, deletion rights, access reviews, and privacy impact assessments must be designed into the platform from inception — not retrofitted after launch.

OWASP ASVS

The web portal, mobile backend APIs, admin console, playback authorization endpoints, device management APIs, and partner-facing APIs must each have defined security verification requirements prior to release. OWASP ASVS governs authentication, session management, authorization, input validation, cryptography, error handling, logging, API abuse protection, and secure file handling across all application surfaces.

NIST Cybersecurity Framework 2.0

The NIST Cybersecurity Framework 2.0 provides the operational control map, organising security posture across six functions: govern, identify, protect, detect, respond, and recover. This structure supports daily security operations, incident response planning, resilience management, and executive reporting — and provides a shared language across security, platform, network, and business stakeholder teams.

9. Security, Privacy, and Compliance Framework

CIS Controls

CIS Controls serve as the implementation checklist for the private-cloud environment, covering asset inventory, secure configuration, vulnerability management, privileged access control, logging, network monitoring, backup management, and incident response. The CIS view complements ISO-style governance by giving infrastructure teams a practical, verifiable hardening baseline.

Video-Specific Security Controls

The following controls apply across live streams, recorded clips, snapshots, thumbnails, metadata, and administrative exports:

- Encrypted transport for all ingest and playback paths
- Private origin access with no direct public exposure
- Signed playback URLs and time-limited tokens
- Strict administrator privilege controls and separation of duties
- Full audit logs for all video access and administrative actions
- Policy-enforced retention and secure, verifiable deletion
- Cryptographic key management with rotation procedures
- Vulnerability scanning, backup and restore testing, and periodic penetration testing

Privacy-by-Design

Privacy controls should be visible in platform behavior. Users and administrators must be able to understand who can access recordings, how long recordings are retained, when deletion occurs, and what actions have been performed on any device or video record. The platform should minimize unnecessary exposure of video content to support staff, provide role-scoped troubleshooting views, and log all privileged access for audit purposes.

Compliance Evidence & Enterprise Readiness

The security architecture must support structured evidence collection. Enterprise customers routinely require compliance reports, penetration test summaries, vulnerability remediation records, backup test results, access review logs, and documented incident response procedures. Designing for evidence from the outset reduces sales friction and positions the platform as mature, controlled, and enterprise-ready.

10. Implementation Duration and Steps

A practical first build of the Video Cloud platform normally takes about 8 to 12 weeks for a minimum viable version and 12 to 20 weeks for a production-hardened version with multi-tenant management, monitoring, usage records, retention automation, and operational dashboards.

The final delivery duration depends on private-cloud readiness, RTSP/ONVIF camera compatibility scope, H.264/H.265 processing requirements, recording retention requirements, security requirements, integration depth, expected concurrency, and acceptance-test standards.

The recommended approach is to deliver in phases: confirm requirements, build the platform foundation, implement the core video workflow, add management and operations capabilities, then complete hardening and acceptance testing.

11. Recommended Implementation Plan

Phase	Duration	Typical Components
Requirements and architecture	1-2 weeks	Scope, protocols, retention, security, tenant model, deployment plan.
Platform foundation	2-3 weeks	Storage, access gateway, database, IAM, logging, monitoring, deployment pipeline.
Core video workflow	4-6 weeks	Ingest, recording, playback, VOD, transcoding, snapshots, watermarking, CDN delivery.
Management and operations	2-4 weeks	Users, permissions, usage records, usage export, dashboards, alerts, audit logs.
Hardening and acceptance	2-4 weeks	Load, failover, security, playback, backup, and operations validation.

Phase 1: Requirements and Architecture
Confirm camera scale assumptions, RTSP/ONVIF scope, H.264/H.265 handling, three-day and 30-day retention rules, mobile/web playback terminals, peak-hour traffic model, security model, compliance baseline, and private-cloud deployment environment.

Phase 2: Platform Foundation
Build private-cloud environments, storage structure, database schema, access gateway, identity/permission baseline, deployment pipeline, logging, monitoring foundation, and secure network segmentation.

Phase 3: Core Video Workflow
Implement ingest, recording, playback, upload/VOD workflow, transcoding, snapshots, watermarking, retention lifecycle, and CDN/origin delivery.

Phase 4: management and operations
Implement user and tenant management, channel/device management, usage records, usage export, dashboards, alerts, audit logs, operation tools, and retention-tier management.

Phase 5: Hardening and Acceptance
Confirm camera scale assumptions, RTSP/ONVIF scope, H.264/H.265 handling, three-day and 30-day retention rules, mobile/web playback terminals, peak-hour traffic model, security model, compliance baseline, and private-cloud deployment environment.

Video Cloud Solution

12. Risk Management

Storage Growth
Mitigate through retention policy, lifecycle tiers, compression settings, deletion automation, and per-tenant quota management.

Ingest Reliability Issues
Mitigate through connection retry logic, access gateway health checks, stream status monitoring, and clear device/channel online-state reporting.

Security and Privacy Exposure
Mitigate through encrypted transport, private storage, signed playback URLs, least-privilege access, audit logs, tenant isolation, and strict administrator controls.

Playback Instability
Mitigate through CDN delivery, tested bitrate settings, protocol compatibility testing, player telemetry, and regional edge strategy.

Operational Complexity
Mitigate through automated deployment, centralized logs, dashboards, alert rules, runbooks, and clear separation between platform services.

Vendor Dependency
Mitigate through standard protocols, portable storage formats, API documentation, infrastructure-as-code, and clear abstraction around cloud-specific media services.

11. Conclusion

The Video Cloud should be understood as strategic infrastructure, not merely a media storage component. Its role is to create a stable and scalable foundation for large-volume video access, recording, playback, management, and governance. When implemented properly, it gives future applications a dependable platform on which to build consumer, enterprise, operational, and analytics-driven services.

The proposed approach emphasizes scale, private-cloud readiness, market-standard storage and CDN patterns, RTSP/ONVIF camera compatibility, H.264/H.265 support, mobile-first playback, strong web administration, and policy-driven retention. These design choices allow the platform to support both immediate deployment needs and long-term growth toward millions of connected cameras.

The strongest differentiators of the platform should be trust and experience. Data security, privacy protection, auditability, and compliance alignment must be visible in the architecture, operations, and product behavior. At the same time, users should experience fast playback, clear service status, reliable recording access, and simple interaction across mobile and web terminals. This combination of engineering depth and user confidence is what positions the Video Cloud as an enterprise-grade solution.

End of paper

SEEING FLOW

Seeingflow Technology Pte. Ltd.
100 Lorong 23 Geylang
#02-02, D'Centennial Building,
Singapore 388398

References

Seeingflow AIOT white paper: https://www.seeingflow.com/documents/sfd-wp-iot-2412.pdf

AWS Video on Demand on AWS implementation guide:
https://docs.aws.amazon.com/solutions/latest/video-on-demand-on-aws/solution-overview.html

AWS Video on Demand on AWS architecture overview:
https://docs.aws.amazon.com/solutions/latest/video-on-demand-on-aws/architecture-overview.html

AWS Live Streaming on AWS implementation guide: https://docs.aws.amazon.com/solutions/latest/live-streaming-on-aws/solution-overview.html

ISO/IEC 27001:2022 information security management systems: https://www.iso.org/standard/27001

ISO/IEC 27017:2015 cloud security controls: https://www.iso.org/standard/43757.html

ISO/IEC 27701 privacy information management: https://www.iso.org/standard/71670.html

NIST Cybersecurity Framework: https://www.nist.gov/cyberframework

OWASP Application Security Verification Standard: https://owasp.org/www-project-application-security-verification-standard/

CIS Critical Security Controls: https://www.cisecurity.org/controls

This white paper is for informational purposes only and is intended to provide an overview of the architecture, capabilities, and design principles of the Video Cloud platform. The information contained herein is provided on an "as is" basis and is subject to change without prior notice.
The content of this document does not constitute a contractual agreement or an offer to sell any product or service. Any features, functionalities, or technical specifications described herein are provided as examples and may vary depending on implementation, deployment environment, customization, and platform updates.

SEEING FLOW

100 Lorong 23 Geylang
#02-02, D'Centennial Building,
Singapore 388398